

Introduction to Statistics with R

Session R02: Correlation

Marvin Schmitt

The example data set

We analyze structures in the **Chicago Face Database** (Ma et al., 2015). Each row refers to a portrait which was rated with respect to different categories by a sample of raters.

```
df = read.csv("R02_notes_dataset.csv")
nrow(df)
```


```
## [1] 597
```

```
colnames(df)
```

```
## [1] "ID"           "Gender"       "Age"         "Afraid"
## [5] "Angry"        "Attractive"   "Babyface"    "Disgusted"
## [9] "Dominant"     "Feminine"     "Happy"       "Masculine"
## [13] "Prototypic"   "Sad"          "Surprised"   "Threatening"
## [17] "Trustworthy"  "Unusual"      "Nose_Width"  "Nose_Length"
## [21] "FaceRoundness" "Noseshape"
```

Scatterplots

```
ggplot(df, aes(x=Masculine, y=Attractive))+  
  geom_point() +  
  ylim(1,5) + xlim(1,5)
```



Digression I: filter

```
df %>% filter(Gender=="M")
```

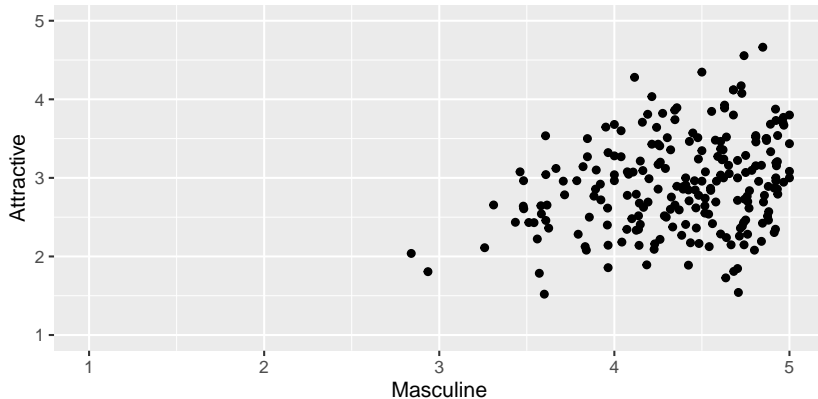
##	ID	Gender	Age	Afraid	Angry	Attractive	Babyface	Disgusted
## 1	58	M	23.80000	1.240000	2.520000	3.040000	2.625000	2.040000
## 2	59	M	26.22222	1.846154	2.888889	2.777778	2.296296	2.296296
## 3	60	M	26.54167	1.708333	1.583333	2.458333	3.041667	1.478261
## 4	61	M	23.91667	1.833333	1.625000	3.041667	3.625000	1.347826
## 5	62	M	34.83333	1.916667	2.130435	2.625000	2.083333	1.708333
## 6	63	M	26.92857	1.785714	2.678571	2.285714	2.428571	2.107143
## 7	64	M	25.91667	1.652174	1.708333	2.958333	2.416667	1.500000
## 8	65	M	23.32143	1.428571	1.535714	2.964286	2.392857	1.750000
## 9	66	M	27.75000	1.592593	1.428571	2.857143	2.678571	1.214286
## 10	67	M	21.04762	2.000000	2.250000	3.809524	3.700000	1.952381
## 11	68	M	26.32143	1.333333	1.851852	3.428571	2.148148	1.321429
## 12	69	M	25.55556	1.925926	3.222222	1.846154	3.307692	2.259259
## 13	70	M	28.07143	1.785714	2.321429	2.142857	2.285714	2.111111
## 14	71	M	29.22222	1.925926	2.259259	3.888889	2.592593	1.814815
## 15	72	M	56.38462	1.884615	1.518519	3.076923	1.851852	2.037037
## 16	73	M	28.48000	1.240000	1.720000	4.120000	2.875000	1.520000
## 17	74	M	43.00000	1.615385	3.076923	2.538462	1.307692	2.538462
## 18	75	M	33.59259	1.250000	1.357143	3.214286	1.964286	1.535714
## 19	76	M	36.90000	1.933333	3.900000	2.517241	1.533333	3.033333
## 20	77	M	23.79310	1.793103	2.862069	2.344828	2.758621	2.357143
## 21	78	M	24.11538	1.884615	2.423077	3.230769	2.846154	2.461538
## 22	79	M	21.92593	1.740741	1.962963	2.222222	2.000000	1.629630
## 23	80	M	28.96000	1.480000	3.040000	1.520000	2.040000	2.800000
## 24	81	M	24.68000	1.920000	2.080000	3.520000	2.840000	2.040000
## 25	82	M	46.81818	1.909091	2.590909	1.809524	2.090909	1.818182
## 26	83	M	41.42308	1.400000	2.076923	2.384615	1.538462	1.384615
## 27	84	M	25.32143	1.642857	2.222222	3.428571	2.642857	1.821429

Digression II: select


```
df %>%  
  select(Attractive)
```

```
##      Attractive  
## 1      4.111111  
## 2      3.111111  
## 3      3.000000  
## 4      3.275862  
## 5      3.172414  
## 6      4.333333  
## 7      2.714286  
## 8      2.137931  
## 9      3.038462  
## 10     4.080000  
## 11     2.615385  
## 12     3.307692  
## 13     2.607143  
## 14     3.000000  
## 15     3.730769  
## 16     2.814815  
## 17     3.266667  
## 18     3.833333  
## 19     4.678571  
## 20     2.928571  
## 21     3.173913  
## 22     3.000000  
## 23     3.689655  
## 24     3.185185  
## 25     3.240000  
## 26     2.024482
```

```
df %>% filter(Gender=="M") %>%  
  ggplot(., aes(x=Masculine, y=Attractive))+  
  geom_point() +  
  ylim(1,5) + xlim(1,5)
```

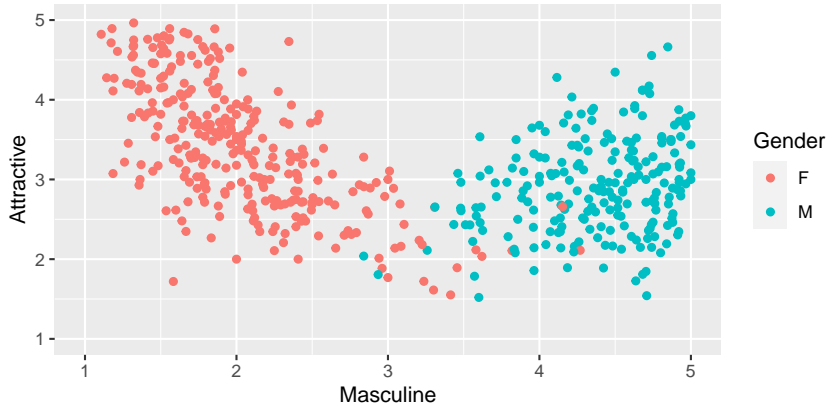


```
df %>% filter(Gender=="F") %>%  
  ggplot(., aes(x=Masculine, y=Attractive))+  
  geom_point() +  
  ylim(1,5) + xlim(1,5)
```




Scatterplots: Group variables

```
ggplot(df, aes(x = Masculine, y = Attractive, color=Gender))+  
  geom_point() +  
  ylim(1, 5) + xlim(1, 5)
```




Scatterplots: Print correlation coefficients

```
ggplot(df, aes(x=Masculine, y=Attractive))+  
  geom_point() +  
  ylim(1, 6) + xlim(1, 5) +  
  stat_cor(p.accuracy = 0.01)
```



```
ggplot(df, aes(x=Masculine, y=Attractive, color=Gender))+  
geom_point() +  
ylim(1, 6.5) + xlim(1, 5) +  
stat_cor(p.accuracy = 0.01)
```



Correlation matrices: Computation

```
df %>%  
  select(Angry, Disgusted, Happy, Sad, Surprised, Attractive, Threatening) %>%  
  cor() %>%  
  round(3)
```

##	Angry	Disgusted	Happy	Sad	Surprised	Attractive	Threatening
## Angry	1.000	0.843	-0.606	0.454	0.082	-0.302	0.834
## Disgusted	0.843	1.000	-0.536	0.572	0.206	-0.296	0.687
## Happy	-0.606	-0.536	1.000	-0.573	0.189	0.463	-0.449
## Sad	0.454	0.572	-0.573	1.000	0.131	-0.323	0.272
## Surprised	0.082	0.206	0.189	0.131	1.000	0.050	0.117
## Attractive	-0.302	-0.296	0.463	-0.323	0.050	1.000	-0.375
## Threatening	0.834	0.687	-0.449	0.272	0.117	-0.375	1.000

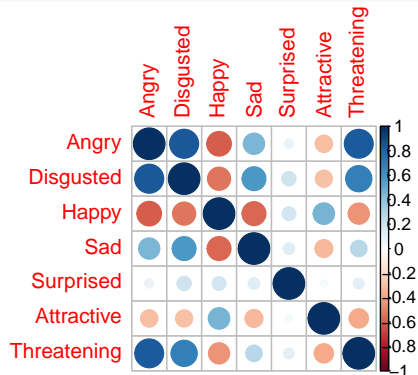
Correlation matrices: Visualization

We can visually encode the entries in a **correlation matrix** while maintaining the matrix structure:

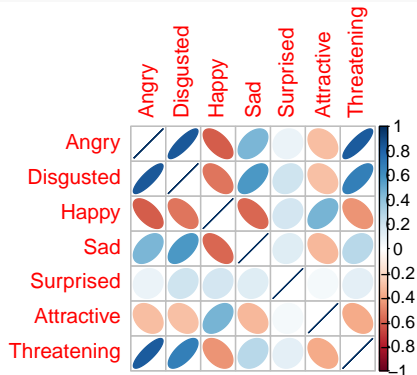
- Color coding
- Circles
- Ellipses
- Correlation coefficients r_{xy}
- ...

We can use the function `corrplot::corrplot` to visualize correlation matrices.

```
df %>%
  select(Angry, Disgusted, Happy, Sad, Surprised, Attractive, Threatening) %>%
  cor() %>%
  corrplot::corrplot()
```



```
df %>%
  select(Angry, Disgusted, Happy, Sad, Surprised, Attractive, Threatening) %>%
  cor() %>%
  corrplot::corrplot(method="ellipse")
```



```
df %>%
  select(Angry, Disgusted, Happy, Sad, Surprised, Attractive, Threatening) %>%
  cor() %>%
  corrplot::corrplot(method="number", number.cex = 0.7)
```



Important arguments for `corrplot::corrplot`:

- `method`: cell content
- `tl.col='black'`: black labels
- `cl.pos='n'`: remove colorbar
- `type`: display entire/lower/upper matrix
- `number.cex`: size of numbers if `method=number`

Testing assumptions

In order to ease interpretation of a correlation coefficient r_{xy} , the variables X, Y should be:

- **metric**,
- **normal**, and
- **homoscedastic**.

We can use the function `shapiro.test(x)` to test `x` for **normality**.

Interpretation:

- if $p \geq .05$: normality assumption not violated
- if $p < .05$: normality assumption violated


```
shapiro.test(df$Trustworthy)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  df$Trustworthy  
## W = 0.99712, p-value = 0.3776
```

```
shapiro.test(df$Attractive)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  df$Attractive  
## W = 0.98357, p-value = 0.000002952
```

The scatterplot allows for a visual inspection of **homoscedasticity**:



Another caveat:

